

ALCF COMPUTATIONAL PERFORMANCE WORKSHOP

MAY 4-6, 2021



CONTAINERS ON THETA AND THETA-GPU AN INTRODUCTION

Romit Maulik, Taylor Childers, Corey Adams

Datascience team

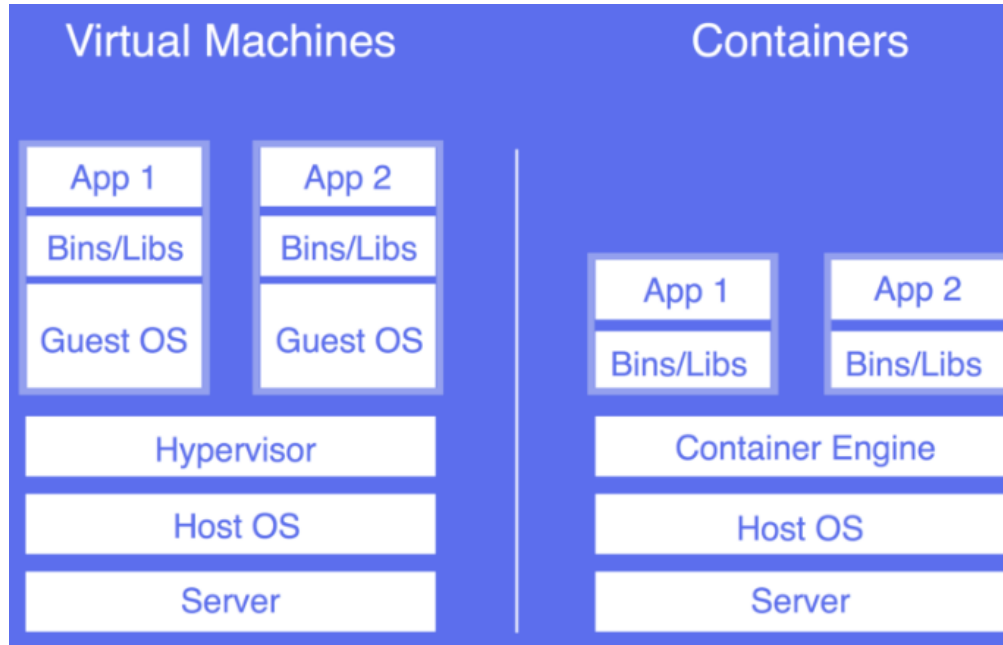
Argonne Leadership Computing Facility

Outline

- Using singularity containers on Theta and Theta-GPU
- Using Docker-hub to build containers and build them using Singularity on Theta (for running MPI apps).
- Using pre-existing singularity images on Theta-GPU for deep learning workloads.
- When in doubt – message us on the CPW portal.

Containers

- Lightweight alternatives to virtual machines and increasingly popular for cloud applications. Examples: Docker, Singularity, Rocket etc.



Singularity is our container of choice for security reasons.

However, Docker is the most widely used container platform.

Dockerfiles and Docker-hub

- Option 1: Build a docker image on their local machine and upload to docker-hub (needs local installation of docker).
- Option 2: Link Github and docker-hub and auto-build using Dockerfile. (does not need local installation but useful for testing anyway).
- We'll show you how docker and docker-hub can be used to port your docker containers to Theta and build Singularity images.
- Due to time constraints – we cannot show you how to install docker on your local machine (but this is not very difficult):

Visit: <https://docs.docker.com/engine/install/> and choose your OS.

Our results/demos today from docker version 20.10.6



source



Dockerfile



submit.sh

An example Dockerfile

```

FROM centos:latest
MAINTAINER Romit rmaulik@anl.gov

WORKDIR /mpich
COPY source/pi.c .

RUN yum update -y
RUN yum groupinstall -y "Development Tools"
RUN yum install -y gcc-c++ wget gcc-gfortran

ENV MPICH_VERSION 3.3
RUN wget http://www.mpich.org/static/downloads/$MPICH_VERSION/mpich-$MPICH_VERSION.tar.gz
RUN tar xf mpich-$MPICH_VERSION.tar.gz --strip-components=1

# disable the addition of the RPATH to compiled executables
# this allows us to override the MPI libraries to use those
# found via LD_LIBRARY_PATH
RUN ./configure --prefix=/mpich/install --disable-wrapper-rpath
RUN make -j 4 install

# add to local environment to build pi.c
ENV PATH $PATH:/mpich/install/bin
ENV LD_LIBRARY_PATH $LD_LIBRARY_PATH:/mpich/install/lib
RUN env | sort
RUN mpicc -o pi -fPIC pi.c

COPY submit.sh .

RUN chmod +x submit.sh

ENTRYPOINT ["/mpich/submit.sh"]

```

Container OS
pulled from docker-hub
repository

Working directory in
container – copy source

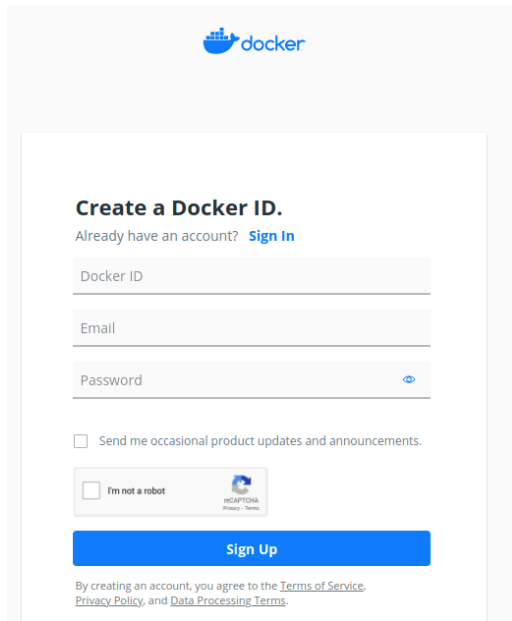
Install appropriate
packages for building
app

Build your app

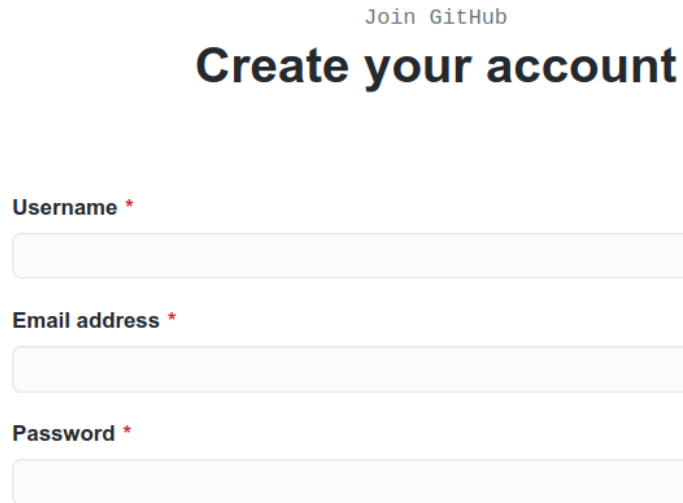
Specify script to run and
copy from local

Linking Github and Docker-hub

- Make sure you have both Github and Docker-hub accounts (sign up with email).



The screenshot shows the Docker website's sign-up page. At the top is the Docker logo. The main heading is 'Create a Docker ID.' Below it is a link 'Already have an account? Sign In'. There are three input fields: 'Docker ID', 'Email', and 'Password' (with an eye icon for toggling visibility). Below the fields is a checkbox for 'Send me occasional product updates and announcements.' and a CAPTCHA section with the text 'I'm not a robot' and a CAPTCHA image. At the bottom is a blue 'Sign Up' button. Below the button is a small disclaimer: 'By creating an account, you agree to the Terms of Service, Privacy Policy, and Data Processing Terms.'

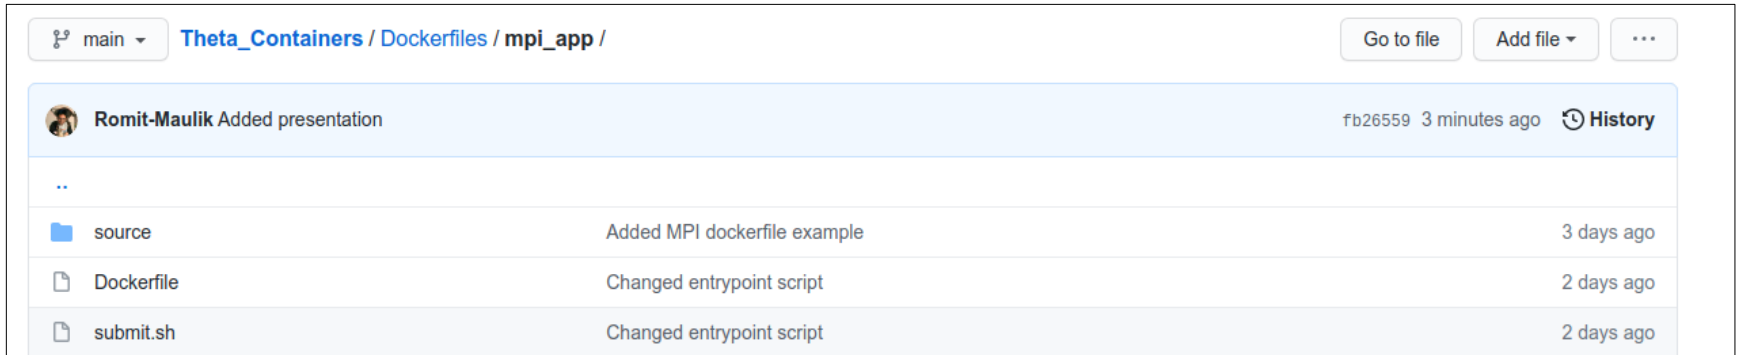


The screenshot shows the GitHub website's sign-up page. At the top is the text 'Join GitHub'. The main heading is 'Create your account'. There are three input fields: 'Username *', 'Email address *', and 'Password *' (with an eye icon for toggling visibility).


Make sure it's at least 15 characters OR at least 8 characters including a number and a lowercase letter. [Learn more.](#)

Linking Github and Docker-hub




- Make sure you have both Github and Docker-hub accounts (sign up with email).
- Push your Dockerfile to a github repository



main ▾ [Theta_Containers](#) / [Dockerfiles](#) / [mpi_app](#) / Go to file Add file ▾ ...

 **Romit-Maulik** Added presentation fb26559 3 minutes ago [History](#)

..

 source	Added MPI dockerfile example	3 days ago
 Dockerfile	Changed entrypoint script	2 days ago
 submit.sh	Changed entrypoint script	2 days ago

Example code at
https://github.com/Romit-Maulik/Theta_Containers/

Linking Github and Docker-hub

- Link that repository with a Docker-hub repository


Create Repository


romitmaulik1 | alcf tutorial1

Using docker/singularity to run containerized apps on ALCF systems

Visibility



Using 0 of 1 private repositories. [Get more](#)

☒ Public  Public repositories appear in Docker Hub search results

☐ Private  Only you can view private repositories

Build Settings (optional)



Autobuild triggers a new build with every git push to your source code repository. [Learn More.](#)





 Connected  Disconnected

[Cancel](#) [Create](#) [Create & Build](#)

Build Settings (optional)

Autobuild triggers a new build with every git push to your source code repository. [Learn More.](#)

 Connected  Disconnected

 Romit-Maulik   Theta_Containers 

BUILD RULES +

The build rules below specify how to build your source into Docker Images.

Source Type	Source	Docker Tag	Dockerfile location	Build Caching
Branch	master	mpl_app	/Dockerfiles/mpl_app	<input checked="" type="checkbox"/>

[View example build rules](#)

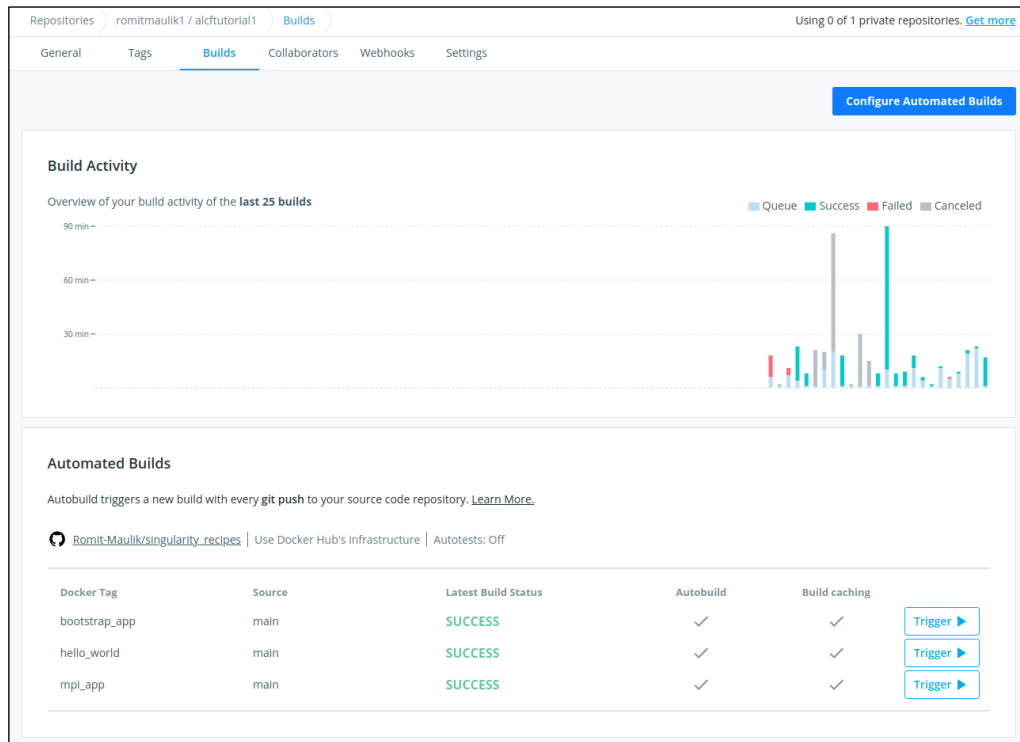
[Cancel](#) [Create](#) [Create & Build](#)

This will be disconnected the first time

Specify branch, image name (tag), path to docker file in repo

Linking Github and Docker-hub

- Building on docker-hub is slow (when compared to your local machine) but once complete you should see something like this



Using “Configure Automated Builds” to edit build names (tags), paths etc.

Recent Builds



[Build in 'main:/Dockerfiles/bootstrap' \(e75e01da\)](#)

Clicking on a recent build will give you logs (useful for debugging fails) and the docker file itself.

Linking Github and Docker-hub

- Building on docker-hub is slow (when compared to your local machine) but once complete you should see something like this

```
BUILD LOGS DOCKERFILE README
Cloning into '.'...
Warning: Permanently added the RSA host key for IP address '140.82.114.4' to the list of known hosts.
Reset branch 'main'
Your branch is up-to-date with 'origin/main'.
KernelVersion: 4.4.0-1060-aws
Components: [[{"Version": "u'19.03.8'", "Name": "u'Engine'", "Details": {"KernelVersion": "u'4.4.0-1060-aws'", "Os": "u'linux'", "BuildTime": "u'2020-03-11T01:24:30.000000000+00:00'", "Arch": "amd64", "BuildTime": "2020-03-11T01:24:30.000000000+00:00", "ApiVersion": "1.40", "Platform": {"Name": "u'Docker Engine - Community'"
Version: 19.03.8
MinAPIVersion: 1.12
GitCommit: afacbb7f0
Os: linux
GoVersion: go1.12.17
Starting build of index.docker.io/ronitmaulik1/alcftutorial1:bootstrap_app...
Step 1/11 : FROM ronitmaulik1/alcftutorial1:npi_app
----> 1d71b9dc4737
Step 2/11 : MAINTAINER Ronit rmaulik@anl.gov
----> Running in 5853cffe9863
Removing intermediate container 5853cffe9863
----> 2c33179ea695
Step 3/11 : WORKDIR /python_app
----> Running in 2c82a6974851
Removing intermediate container 2c82a6974851
----> b563dd74d077
```

Using “Configure Automated Builds” to edit build names (tags), paths etc.

Recent Builds



[Build in 'main:/Dockerfiles/bootstrap' \(e75e01da\)](#)

Clicking on a recent build will give you logs (useful for debugging fails) and the docker file itself.

Linking Github and Docker-hub

- Building on docker-hub is slow (when compared to your local machine) but once complete you should see something like this

```
BUILD LOGS DOCKERFILE README

FROM centos:latest
MAINTAINER Romit rmaulik@anl.gov

WORKDIR /mpich
COPY source/pi.c .

RUN yum update -y
RUN yum groupinstall -y "Development Tools"
RUN yum install -y gcc-c++ wget gcc-gfortran

ENV MPICH_VERSION 3.3
RUN wget http://www.mpich.org/static/downloads/$MPICH_VERSION/mpich-$MPICH_VERSION.tar.gz
RUN tar xf mpich-$MPICH_VERSION.tar.gz --strip-components=1

# disable the addition of the RPATH to compiled executables
# this allows us to override the MPI libraries to use those
# found via LD_LIBRARY_PATH
RUN ./configure --prefix=/mpich/install --disable-wrapper-rpath
RUN make -j 4 install

# add to local environment to build pi.c
ENV PATH $PATH:/mpich/install/bin
ENV LD_LIBRARY_PATH $LD_LIBRARY_PATH:/mpich/install/lib
RUN env | sort
RUN mpicc -o pi -fPIC pi.c

COPY submit.sh .

RUN chmod +x submit.sh

ENTRYPOINT ["/mpich/submit.sh"]
```

Using “Configure Automated Builds” to edit build names (tags), paths etc.

Recent Builds



[Build in 'main:/Dockerfiles/bootstrap' \(e75e01da\)](#)

Clicking on a recent build will give you logs (useful for debugging fails) and the docker file itself.

Pulling image to Theta

- Once build is successful on Docker-hub – pull to Theta using singularity (version 3.6.4-1)

- Command:

```
singularity build mpi_app.img docker://romitmaulik1/alcftutorial1:mpi_app
```

- Explore the contents of your image with:

```
singularity shell mpi_app.img
```

- You can find your source code, downloaded libraries, submit scripts in

```
ls /mpich/
```

- Your entrypoint script will look like

```
1  #!/bin/bash
2  # submit.sh
3
4  echo "running MPI app"
5  cd /mpich/
6  ./pi
7  |
```

Interface MPI apps with Theta

- We need to swap out binaries with Cray in the submission script

```
2 #!/bin/bash
3 #COBALT -t 30
4 #COBALT -q debug-cache-quad
5 #COBALT -n 2
6
7 RANKS_PER_NODE=4
8
9 # pass container as first argument to script
10 CONTAINER=$1
11
12 # Use Cray's Application Binary Independent MPI build
13 module swap cray-mpich cray-mpich-abi
14
15 # Output current modules being used (for debugging)
16 module list
17
18 # Only needed when interactive debugging
19 #module swap PrgEnv-intel PrgEnv-cray; module swap PrgEnv-cray PrgEnv-intel
20
21 export ADDITIONAL_PATHS="/opt/cray/diag/lib:/opt/cray/ugni/default/lib64:/opt/cray/udreg/default/lib64:/opt
22
23 # in order to pass environment variables to a Singularity container create the
24 # variable with the SINGULARITYENV prefix
25 export SINGULARITYENV_LD_LIBRARY_PATH="$CRAY_LD_LIBRARY_PATH:$LD_LIBRARY_PATH:$ADDITIONAL_PATHS"
26
27 # print to log file for debug
28 echo $SINGULARITYENV_LD_LIBRARY_PATH
29
30 TOTAL_RANKS=$(( $COBALT_JOBSIZE * $RANKS_PER_NODE ))
31
32 # this simply runs the command 'ldd /myapp/pi' inside the container and should
33 # show that the app is running againsts the host machines Cray libmpi.so not the
34 # one inside the container
35 BINDINGS="-B /opt -B /etc/alternatives"
36 aprun -n 1 -N 1 singularity exec $BINDINGS $CONTAINER bash -c "echo \$LD_LIBRARY_PATH"
37 aprun -n 1 -N 1 singularity exec $BINDINGS $CONTAINER bash -c "ldd /myapp/pi"
38
39 # run my container like an application, which will run '/myapp/pi'
40 aprun -n $TOTAL_RANKS -N $RANKS_PER_NODE singularity run $BINDINGS $CONTAINER
```

Use cray-mpich-abi for compatibility with your build (remember to check compatibility of versions of mpich and cray-mpich)

Pass paths to CRAY objects to your container using SINGULARITYENV_LD_LIBRARY_PATH

Run using aprun

What is finally observed

```
linux-vdso.so.1 (0x00002aaaaad7000)
libmpi.so.12 => /opt/cray/pe/mpt/7.7.14/gni/mpich-intel-abi/16.0/lib/libmpi.so.12 (0x00002aaaaacd6000)
libc.so.6 => /lib64/libc.so.6 (0x00002aaaab299000)
libxpmem.so.0 => /opt/cray/xpmem/default/lib64/libxpmem.so.0 (0x00002aaaab65c000)
librt.so.1 => /lib64/librt.so.1 (0x00002aaaab85f000)
libugni.so.0 => /opt/cray/ugni/default/lib64/libugni.so.0 (0x00002aaaaba67000)
libudreg.so.0 => /opt/cray/udreg/default/lib64/libudreg.so.0 (0x00002aaaabceb000)
libpthread.so.0 => /lib64/libpthread.so.0 (0x00002aaaabef5000)
libpmi.so.0 => /opt/cray/pe/pmi/5.0.16/lib64/libpmi.so.0 (0x00002aaaac115000)
libifport.so.5 => /opt/intel/compilers_and_libraries_2020.0.166/linux/compiler/lib/intel64/libifport.so.5 (0x00002aaaac35e000)
libifcore.so.5 => /opt/intel/compilers_and_libraries_2020.0.166/linux/compiler/lib/intel64/libifcore.so.5 (0x00002aaaac58c000)
libimf.so => /opt/intel/compilers_and_libraries_2020.0.166/linux/compiler/lib/intel64/libimf.so (0x00002aaaac8f0000)
libsvml.so => /opt/intel/compilers_and_libraries_2020.0.166/linux/compiler/lib/intel64/libsvml.so (0x00002aaaace8e000)
libm.so.6 => /lib64/libm.so.6 (0x00002aaaae815000)
libintlc.so.5 => /opt/intel/compilers_and_libraries_2020.0.166/linux/compiler/lib/intel64/libintlc.so.5 (0x00002aaaaeb97000)
/lib64/ld-linux-x86-64.so.2 (0x00002aaaaaab0000)
libgcc_s.so.1 => /lib64/libgcc_s.so.1 (0x00002aaaaee0e000)
libdl.so.2 => /lib64/libdl.so.2 (0x00002aaaaf026000)
Application 22973002 resources: utime ~1s, stime ~2s, Rss ~37072, inblocks ~18068, outblocks ~0
running MPI app
running MPI app
running MPI app
running MPI app
running MPI app
running MPI app
running MPI app
running MPI app
worker 3 of 8
worker 2 of 8
worker 1 of 8
worker 0 of 8
worker 6 of 8
worker 7 of 8
worker 4 of 8
worker 5 of 8
pi is approximately 3.1417259869152532, Error is 0.0001333333254601
Application 22973003 resources: utime ~10s, stime ~8s, Rss ~36056, inblocks ~41708, outblocks ~0
```

Bootstrapping

- When possible reuse previous builds in opensource repositories

```
1 # Bootstrap our previous image
2 FROM romitmaulik1/alcftutorial1:mpi_app
3 MAINTAINER Romit rmaulik@anl.gov
4
5 WORKDIR /python_app
6 COPY source/hello_world.py .|
7
8 RUN yum update -y && yum -y install epel-release
9 RUN yum -y install https://repo.ius.io/ius-release-el7.rpm https://dl.fedoraproject.org/pub/epel/
   epel-release-latest-7.noarch.rpm
10
11 RUN yum -y makecache && yum -y install python3 python3-pip python3-devel && yum clean all
12
13 RUN pip3 install numpy tensorflow matplotlib
14
15 COPY submit.sh .
16
17 RUN chmod +x submit.sh
18
19 ENTRYPOINT ["/python_app/submit.sh"]
20
```

We add a python application into this pre-existing image and build a new one

Install some common datascience packages

Bootstrapping

- When possible reuse previous builds in opensource repositories

```
rmaulik@thetalogin5:~/singularity_tutorial/bootstrap> singularity run bootstrap_app.img
running bootstrap app
/usr/local/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:516: FutureWarning: Passing
, it will be understood as (type, (1,)) / '(1,)type'.
  np_qint8 = np.dtype [("qint8", np.int8, 1)]
/usr/local/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:517: FutureWarning: Passing
, it will be understood as (type, (1,)) / '(1,)type'.
  np_quint8 = np.dtype [("quint8", np.uint8, 1)]
/usr/local/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:518: FutureWarning: Passing
, it will be understood as (type, (1,)) / '(1,)type'.
  np_qint16 = np.dtype [("qint16", np.int16, 1)]
/usr/local/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:519: FutureWarning: Passing
, it will be understood as (type, (1,)) / '(1,)type'.
  np_quint16 = np.dtype [("quint16", np.uint16, 1)]
/usr/local/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:520: FutureWarning: Passing
, it will be understood as (type, (1,)) / '(1,)type'.
  np_qint32 = np.dtype [("qint32", np.int32, 1)]
/usr/local/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:525: FutureWarning: Passing
, it will be understood as (type, (1,)) / '(1,)type'.
  np_resource = np.dtype [("resource", np.ubyte, 1)]
/usr/local/lib/python3.6/site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:541: FutureWarning:
f numpy, it will be understood as (type, (1,)) / '(1,)type'.
  np_qint8 = np.dtype [("qint8", np.int8, 1)]
/usr/local/lib/python3.6/site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:542: FutureWarning:
f numpy, it will be understood as (type, (1,)) / '(1,)type'.
  np_quint8 = np.dtype [("quint8", np.uint8, 1)]
/usr/local/lib/python3.6/site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:543: FutureWarning:
f numpy, it will be understood as (type, (1,)) / '(1,)type'.
  np_qint16 = np.dtype [("qint16", np.int16, 1)]
/usr/local/lib/python3.6/site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:544: FutureWarning:
f numpy, it will be understood as (type, (1,)) / '(1,)type'.
  np_quint16 = np.dtype [("quint16", np.uint16, 1)]
/usr/local/lib/python3.6/site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:545: FutureWarning:
f numpy, it will be understood as (type, (1,)) / '(1,)type'.
  np_qint32 = np.dtype [("qint32", np.int32, 1)]
/usr/local/lib/python3.6/site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:550: FutureWarning:
f numpy, it will be understood as (type, (1,)) / '(1,)type'.
  np_resource = np.dtype [("resource", np.ubyte, 1)]
Hello world
Numpy version: 1.19.5
tensorflow version: 1.19.5
```

```
1 import numpy as np
2 import tensorflow as tf
3
4 print('Hello world')
5 print('Numpy version:', np.__version__)
6 print('tensorflow version:', np.__version__)
```


ThetaGPU

- There are several containers on ThetaGPU that will help you get started with deep learning experiments that can efficiently use the A100 GPUs.
- Documentation: <https://argonne-lcf.github.io/ThetaGPU-Docs/>
- The different optimized containers for DL are available at:
</lus/theta-fs0/projects/datascience/thetaGPU/containers/>
- To leverage these containers – you must either perform the following in a shell script that acquires a compute node or do this interactively.
- Further details may be found in:
https://github.com/argonne-lcf/sdl_ai_workshop/tree/master/05_Simulation_ML/ThetaGPU

ThetaGPU

- For interactive/debug DL, grab an interactive node with:

```
qsub -n 1 -q training -A project_name -l -t 1:00:00
```

- Activate the singularity container:

```
singularity exec -B /lus:/lus --nv /lus/theta-fs0/software/thetagpu/nvidia-containers/tensorflow2/tf2_20.08-py3.simg bash
```

- If your application needs new packages, setup access to the internet:

```
export http_proxy=http://theta-proxy.tmi.alcf.anl.gov:3128
```

```
export https_proxy=https://theta-proxy.tmi.alcf.anl.gov:3128
```

- Create a virtual environment and install your packages for example:

```
python -m pip install --user virtualenv
```

```
export VENV_LOCATION=/home/$USER/THETAGPU_TF_ENV # Add your path here
```

```
python -m virtualenv --system-site-packages $VENV_LOCATION
```

```
source $VENV_LOCATION/bin/activate
```

```
pip install sklearn
```

ThetaGPU

- Now that packages are installed – run your DL script like you would on your local machine or construct a script for submission to the queue:

```
#!/bin/bash
```

```
#COBALT -n 1
```

```
#COBALT -t 00:10:00
```

```
#COBALT -q training
```

```
#COBALT -A project_name
```

```
CONTAINER=/lus/theta-fs0/software/thetagpu/nvidia-containers/tensorflow2/tf2_20.08-py3.simg
```

```
SCRIPT=/path/to/project_location/queue_submission.sh
```

```
echo "Running Cobalt Job $COBALT_JOBID."
```

```
mpirun -n 1 -npnode 1 -hostfile $COBALT_NODEFILE singularity run --nv -B /lus:/lus  
$CONTAINER $SCRIPT
```

ThetaGPU

- Where `queue_submission.sh` is:

```
#!/bin/bash
```

```
cd /path/to/project_location
```

```
export VENV_LOCATION=/home/rmaulik/THETAGPU_TF_ENV
```

```
source $VENV_LOCATION/bin/activate
```

```
python myscript.py
```

- You may also use Docker-hub to build images on the fly on ThetaGPU (after exporting proxies) – although we recommend using what we just discussed to add your libraries to a virtual environment.

Fin!

- I encourage you to try these tutorials and get back to us with any questions.